

# Assignment 1 – Deep Neural Networks (DNN)

**Programming Assignment: Comparing Linear Models and MLPs,** Max Marks: 10

Course: Deep Neural Networks, Type: Individual Programming

Submission: .ipynb + HTML version ( Download HTML version of the ipynb file)

## 1. Overview

**What You Will Do -** Select a real-world dataset ( $\geq 500$  samples,  $\geq 10 < 50$  features), implement a baseline linear model from scratch (no sklearn models), implement an MLP from scratch, compare performance on test data, analyze results.

**Learning Objectives -** Understand regression/classification fundamentals, implement gradient descent, build neural networks without high-level libraries, evaluate and compare models, analyze computational trade-offs.

## 2. Dataset Requirements

Source: UCI ML Repository, Kaggle, OpenML, Data.gov

Samples:  $\geq 500$

Features:  $\geq 10$  (excluding target)

Data type: Numeric and/or categorical

Task: **Regression or Classification (Multiclass)**

**Not Allowed Datasets-** Image datasets (MNIST, CIFAR), text/NLP datasets, audio datasets, preprocessed datasets from Papers with Code, toy datasets with  $< 500$  samples (Iris, Wine, Diabetes).

## 3. Implementation Requirements

### Must Implement From Scratch

Regression: Linear Regression + MLP (with backprop)

Binary Classification: Logistic Regression + MLP

Multi-class: Softmax Regression + MLP

**Allowed Libraries** - NumPy (arrays), Pandas (data handling), Matplotlib/Seaborn (plots), sklearn ONLY for train\_test\_split, StandardScaler/MinMaxScaler, LabelEncoder/OneHotEncoder.

**Prohibited** - sklearn.linear\_model, sklearn.neural\_network, TensorFlow, Keras, PyTorch, JAX, any high-level ML library.

Penalty for using prohibited libraries: -5 marks.

## 4. Assignment Structure

**4.1 Dataset Selection (1 Mark)** - Load dataset, describe it, state regression/classification, justify primary evaluation metric important for your dataset.

**4.2 Data Preprocessing (1 Mark)** - Train-test split (70-30/80-20/ 90-10), handle missing values, encode categorical variables, scale features, etc.

**4.3 Baseline Model (2 Marks)** - Implement from scratch, initialize weights, forward pass, compute loss, compute gradients, update weights ( $w = w - lr \times \text{grad}$ ), store loss at every iteration (loss\_history required), implement predict() function which will return prediction.

**4.4 Multi-Layer Perceptron (4 Marks)** - At least 1 hidden layer, ReLU for hidden layers, appropriate output activation, full forward propagation, backpropagation using chain rule, gradient descent updates, loss\_history required, architecture format: [input, hidden1, hidden2, ..., output].

Mandatory functions inside MLP class:

- init() ← for architecture initialisation
- initialize\_parameters() ← (initialize W and b for each layer)
- forward\_propagation() ← # Compute activations through all layers
- backward\_propagation() ← # Compute gradients using chain rule
- fit() ← # Training loop with forward / backward passes
- predict() ← # Return predictions

Implement function get\_assignment\_results().

Return: dataset\_name, n\_samples, n\_features, problem\_type, primary\_metric, baseline\_model metrics (including test performance, training time), mlp\_model metrics (including architecture, test performance, training time).

Missing function: -2 marks.

## **4.5 Evaluation & Comparison (2 Marks)**

Regression: MSE, RMSE, MAE, R<sup>2</sup>

Classification: Accuracy, Precision, Recall, F1, etc

Required plots: training loss curves (baseline vs MLP), performance comparison bar chart, optional domain-specific plots.

Analysis (< 200 words): which model performed better, by how much, why, computational cost difference, challenges/surprises.

### **Debugging Checklist**

If model not working:

1. Check data shapes at each step, 2. Verify gradient computation
3. Check for NaN/Inf values, 4. Try learning rates: [0.001, 0.01, 0.1]
5. Ensure loss computed correctly, 6. Verify update:  $w = w - lr * \text{grad}$
7. Check activation functions, 8. Print intermediate values

## **5. Evaluation Criteria (10 Marks)**

Dataset: 1, Preprocessing: 1, Baseline (correct implementation, decreasing loss): 2

MLP (valid architecture, forward/backprop, decreasing loss): 4

Comparison (metrics, plots, < 200 words analysis): 2

Deductions: Dataset too small: -1, Loss not decreasing: -1 per model, Analysis  $\geq$  200 words: -0.5, Prohibited libraries: -5

## **6. Submission Instructions**

File name format: YourStudentID\_assignment1.ipynb,

Example: 2025AA05005\_assignment1.ipynb. Upload to LMS, verify successful submission.

## **7. Academic Integrity**

Allowed: Clarifications, conceptual help, debugging assistance.

Not allowed: Sharing code, copying implementations, using AI for full solutions, collaboration. Individual assignment.

---

Start early, scale features, track loss carefully, use vectorized NumPy, test frequently, ensure both models converge.